

# TASK - 3

## Importing necessary libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
%matplotlib inline
```

## Reading the dataset

```
In [2]: df=pd.read_csv("bank-additional.csv",delimiter=';')
# delimiter is used to seperate the columns

df.rename(columns={'y':'deposit'}, inplace=True)
df.head()
```

```
Out[2]:   age    job marital education default housing loan contact month day_of_week ... campaign pdays previous
0   30  blue-collar married basic.9y no yes no cellular may fri ... 2 999
1   39    services single high.school no no no telephone may fri ... 4 999
2   25    services married high.school no yes no telephone jun wed ... 1 999
3   38    services married basic.9y no unknown unknown telephone jun fri ... 3 999
4   47   admin. married university.degree no yes no cellular nov mon ... 1 999
```

5 rows × 21 columns

## Some information about the dataset

```
In [3]: df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4119 entries, 0 to 4118
Data columns (total 21 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age         4119 non-null   int64  
 1   job          4119 non-null   object  
 2   marital     4119 non-null   object  
 3   education   4119 non-null   object  
 4   default     4119 non-null   object  
 5   housing     4119 non-null   object  
 6   loan         4119 non-null   object  
 7   contact     4119 non-null   object  
 8   month        4119 non-null   object  
 9   day_of_week 4119 non-null   object  
 10  duration    4119 non-null   int64  
 11  campaign    4119 non-null   int64  
 12  pdays       4119 non-null   int64  
 13  previous    4119 non-null   int64  
 14  poutcome    4119 non-null   object  
 15  emp.var.rate 4119 non-null   float64 
 16  cons.price.idx 4119 non-null   float64 
 17  cons.conf.idx 4119 non-null   float64 
 18  euribor3m   4119 non-null   float64 
 19  nr.employed 4119 non-null   float64 
 20  deposit     4119 non-null   object  
dtypes: float64(5), int64(5), object(11)
memory usage: 675.9+ KB
```

```
In [4]: df.describe()
```

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.idx	euribor3m
<b>count</b>	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000	4119.000000
<b>mean</b>	40.113620	256.788055	2.537266	960.422190	0.190337	0.084972	93.579704	-40.499102	3.621356
<b>std</b>	10.313362	254.703736	2.568159	191.922786	0.541788	1.563114	0.579349	4.594578	1.733591
<b>min</b>	18.000000	0.000000	1.000000	0.000000	0.000000	-3.400000	92.201000	-50.800000	0.635000
<b>25%</b>	32.000000	103.000000	1.000000	999.000000	0.000000	-1.800000	93.075000	-42.700000	1.334000
<b>50%</b>	38.000000	181.000000	2.000000	999.000000	0.000000	1.100000	93.749000	-41.800000	4.857000
<b>75%</b>	47.000000	317.000000	3.000000	999.000000	0.000000	1.400000	93.994000	-36.400000	4.961000
<b>max</b>	88.000000	3643.000000	35.000000	999.000000	6.000000	1.400000	94.767000	-26.900000	5.045000

In [5]: df.shape

Out[5]: (4119, 21)

In [6]: df.columns

```
Out[6]: Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
   'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
   'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
   'cons.conf.idx', 'euribor3m', 'nr.employed', 'deposit'],
  dtype='object')
```

## Checking for the datatypes

In [7]: df.dtypes

```
Out[7]: age          int64
job           object
marital        object
education      object
default         object
housing         object
loan            object
contact         object
month           object
day_of_week    object
duration        int64
campaign        int64
pdays           int64
previous        int64
poutcome        object
emp.var.rate    float64
cons.price.idx float64
cons.conf.idx  float64
euribor3m       float64
nr.employed     float64
deposit          object
dtype: object
```

In [8]: df.dtypes.value\_counts()

```
Out[8]: object    11
int64     5
float64   5
Name: count, dtype: int64
```

## Checking for the duplicate values

In [9]: df.duplicated().sum()

Out[9]: np.int64(0)

## Checking for null values

In [10]: df.isna().sum()

```
Out[10]: age          0
          job          0
          marital      0
          education    0
          default      0
          housing      0
          loan          0
          contact       0
          month         0
          day_of_week   0
          duration      0
          campaign      0
          pdays         0
          previous      0
          poutcome      0
          emp.var.rate  0
          cons.price.idx 0
          cons.conf.idx  0
          euribor3m     0
          nr.employed   0
          deposit        0
          dtype: int64
```

```
In [11]: cat_cols = df.select_dtypes(include='object').columns
print(cat_cols)
```

```
num_cols = df.select_dtypes(exclude='object').columns
print(num_cols)
```

```
Index(['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
       'month', 'day_of_week', 'poutcome', 'deposit'],
      dtype='object')
```

```
Index(['age', 'duration', 'campaign', 'pdays', 'previous', 'emp.var.rate',
       'cons.price.idx', 'cons.conf.idx', 'euribor3m', 'nr.employed'],
      dtype='object')
```

```
In [12]: df.describe()
```

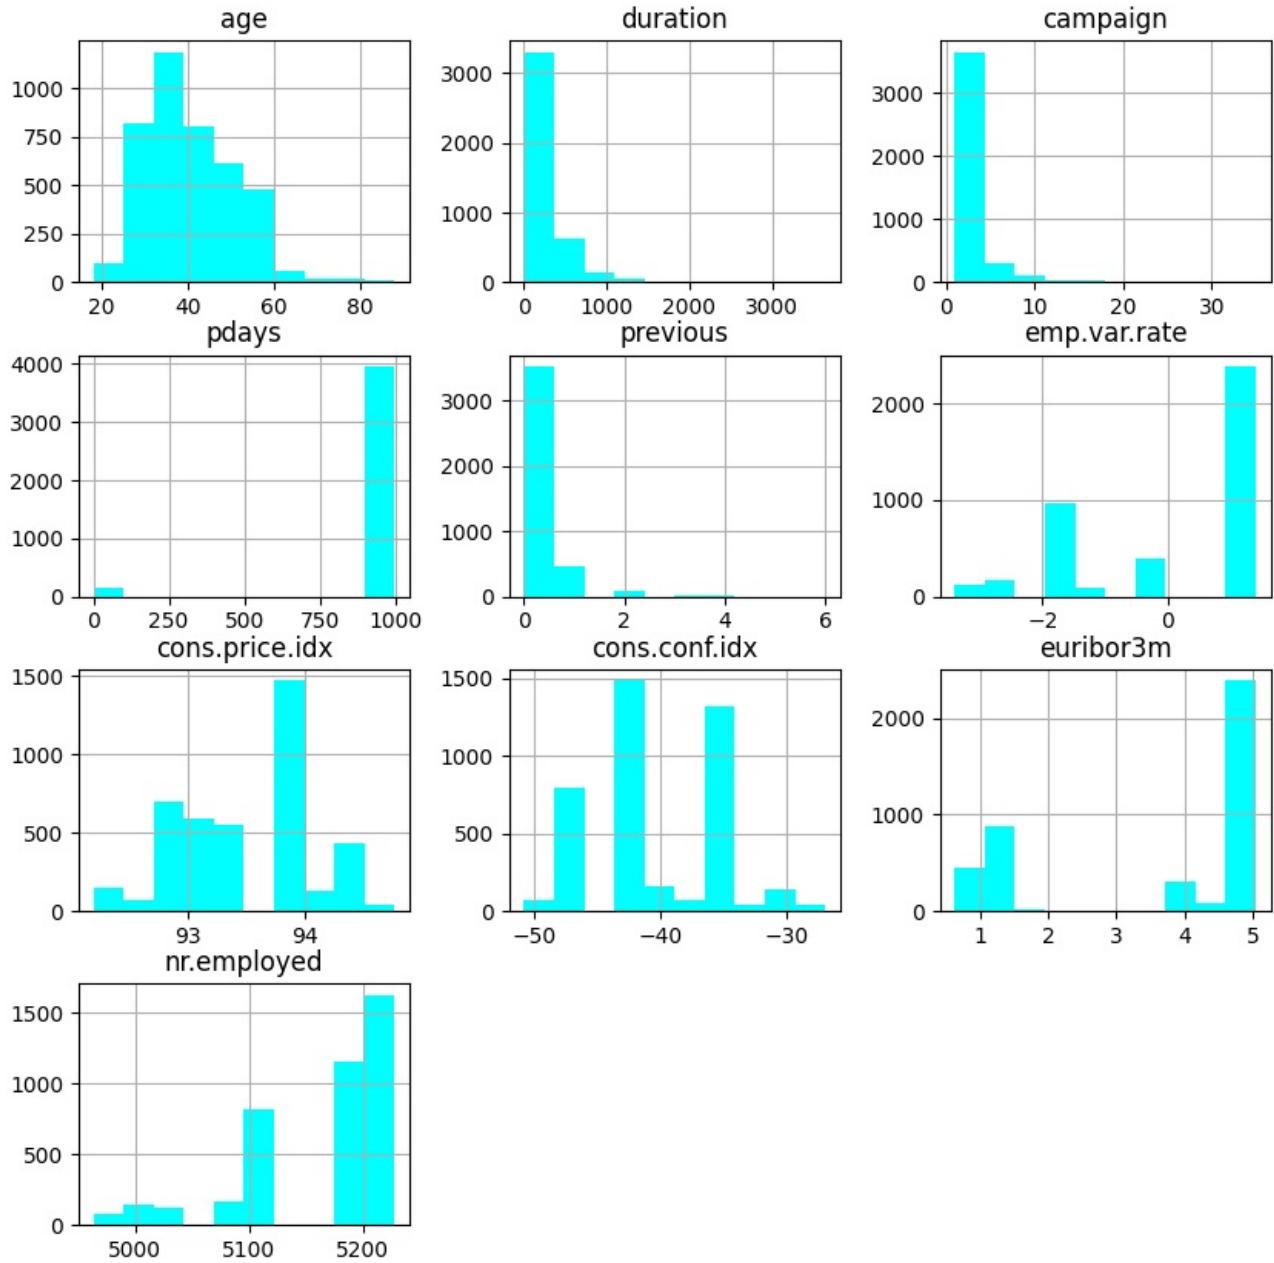
```
Out[12]:      age  duration  campaign  pdays  previous  emp.var.rate  cons.price.idx  cons.conf.idx  euribor3m
count  4119.000000  4119.000000  4119.000000  4119.000000  4119.000000  4119.000000  4119.000000  4119.000000  4119.000000
mean   40.113620  256.788055  2.537266  960.422190  0.190337  0.084972  93.579704  -40.499102  3.621356
std    10.313362  254.703736  2.568159  191.922786  0.541788  1.563114  0.579349  4.594578  1.733591
min    18.000000  0.000000  1.000000  0.000000  0.000000  -3.400000  92.201000  -50.800000  0.635000
25%   32.000000  103.000000  1.000000  999.000000  0.000000  -1.800000  93.075000  -42.700000  1.334000
50%   38.000000  181.000000  2.000000  999.000000  0.000000  1.100000  93.749000  -41.800000  4.857000
75%   47.000000  317.000000  3.000000  999.000000  0.000000  1.400000  93.994000  -36.400000  4.961000
max   88.000000  3643.000000 35.000000  999.000000  6.000000  1.400000  94.767000  -26.900000  5.045000
```

```
In [13]: df.describe(include='object')
```

```
Out[13]:      job  marital  education  default  housing  loan  contact  month  day_of_week  poutcome  deposit
count   4119    4119      4119     4119     4119    4119     4119    4119      4119     4119     4119
unique    12       4        8       3       3       3       2      10        5       3       2
top    admin. married university.degree    no     yes     no cellular    may    thu nonexistent    no
freq  1012   2509      1264     3315    2175    3349    2652    1378      860     3523    3668
```

## Visualisation

```
In [14]: df.hist(figsize=(10,10),color="#00FFFF")
plt.show()
```

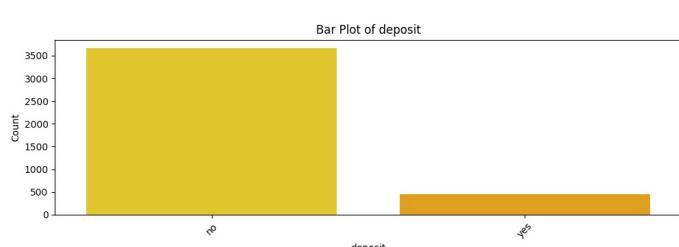
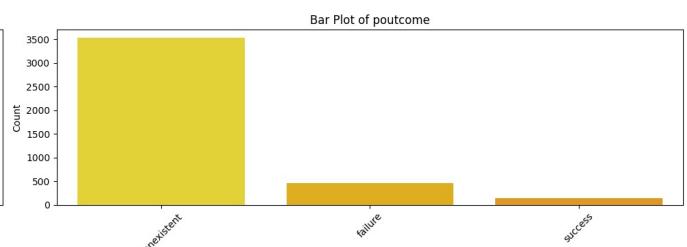
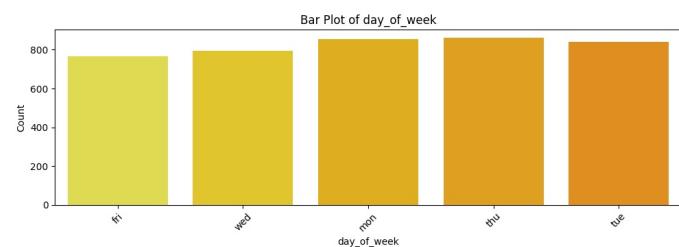
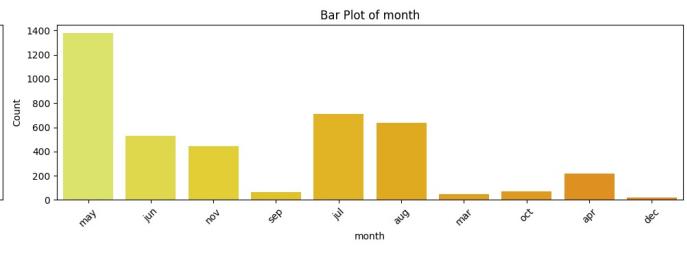
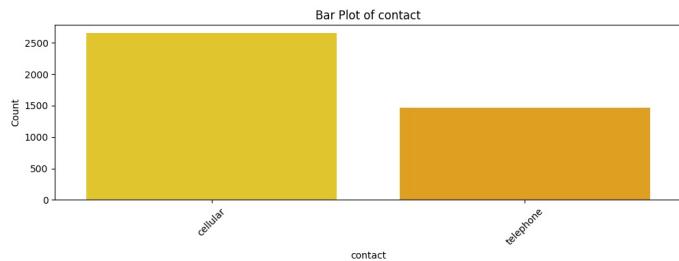
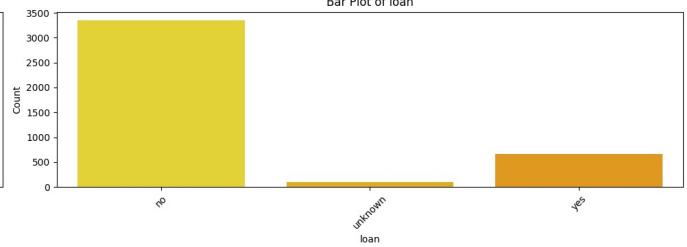
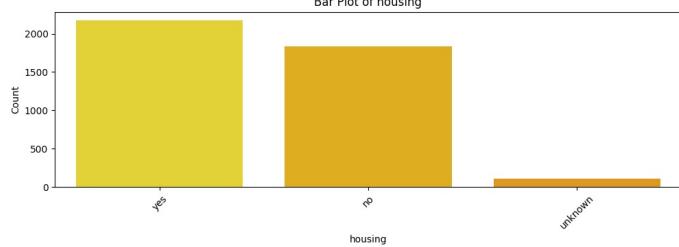
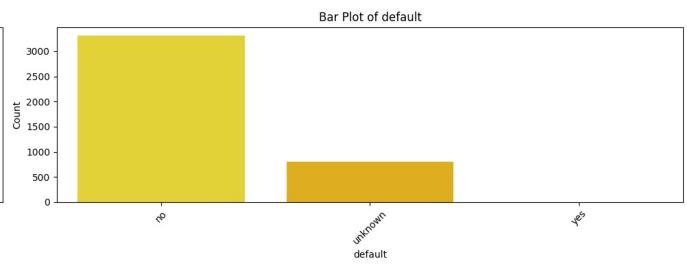
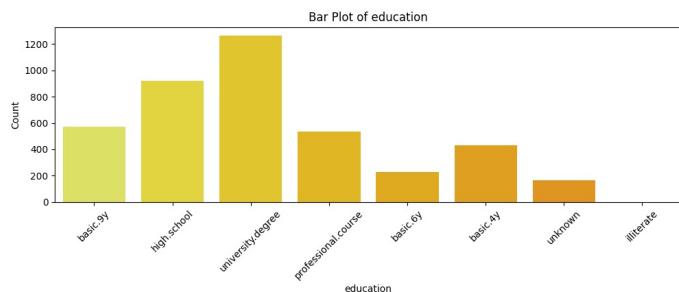
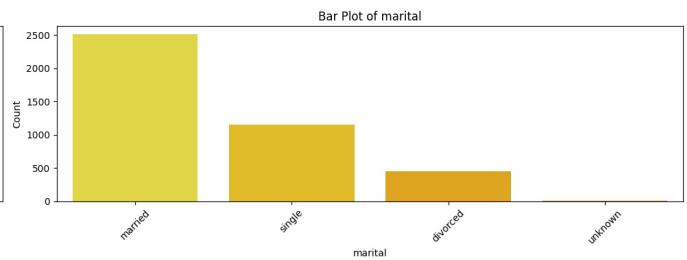
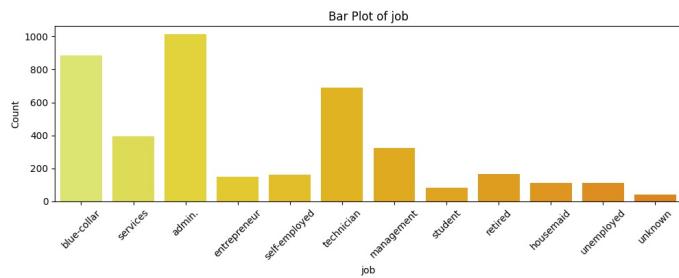


```
In [15]: # Calculate the number of rows and columns for subplots
num_plots = len(cat_cols)
num_rows = (num_plots + 1) // 2 # Add 1 and divide by 2 to round up for odd numbers
num_cols = 2

# Create a new figure
plt.figure(figsize=(20, 25)) # Adjust the figure size as needed

# Loop through each feature and create a countplot
for i, feature in enumerate(cat_cols, 1):
    plt.subplot(num_rows, num_cols, i)
    sns.countplot(x=feature, data=df, palette='Wistia')
    plt.title(f'Bar Plot of {feature}')
    plt.xlabel(feature)
    plt.ylabel('Count')
    plt.xticks(rotation=45)

# Adjust layout to prevent overlap of subplots
plt.tight_layout()
plt.show()
```



```
In [16]: high_corr_cols = ['emp.var.rate', 'euribor3m', 'nr.employed']
```

```
In [17]: df1 = df.copy()
df1.columns
```

```
Out[17]: Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
       'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
       'previous', 'poutcome', 'emp.var.rate', 'cons.price.idx',
       'cons.conf.idx', 'euribor3m', 'nr.employed', 'deposit'],
      dtype='object')
```

```
In [18]: df1.drop(high_corr_cols,inplace=True,axis=1) # axis=1 indicates columns
df1.columns
```

```
Out[18]: Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',  
       'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',  
       'previous', 'poutcome', 'cons.price.idx', 'cons.conf.idx', 'deposit'],  
      dtype='object')
```

```
In [19]: %pip install scikit-learn
```

```
Requirement already satisfied: scikit-learn in c:\users\komal\appdata\local\programs\python\python314\lib\site-packages (1.8.0)  
Requirement already satisfied: numpy>=1.24.1 in c:\users\komal\appdata\local\programs\python\python314\lib\site-packages (from scikit-learn) (2.4.0)  
Requirement already satisfied: scipy>=1.10.0 in c:\users\komal\appdata\local\programs\python\python314\lib\site-packages (from scikit-learn) (1.16.3)  
Requirement already satisfied: joblib>=1.3.0 in c:\users\komal\appdata\local\programs\python\python314\lib\site-packages (from scikit-learn) (1.5.3)  
Requirement already satisfied: threadpoolctl>=3.2.0 in c:\users\komal\appdata\local\programs\python\python314\lib\site-packages (from scikit-learn) (3.6.0)  
Note: you may need to restart the kernel to use updated packages.
```

```
In [20]: from sklearn.preprocessing import LabelEncoder  
lb = LabelEncoder()  
df_encoded = df1.apply(lb.fit_transform)  
df_encoded
```

```
Out[20]:   age  job  marital  education  default  housing  loan  contact  month  day_of_week  duration  campaign  pdays  previous  p  
0    12     1        1         1         2         0         2         0         0         6         0        474         1        20         0  
1    21     7        2         2         3         0         0         0         1         6         0        343         3        20         0  
2     7     7        1         1         3         0         2         0         1         4         4        224         0        20         0  
3    20     7        1         1         2         0         1         1         1         4         0        14         2        20         0  
4    29     0        1         1         6         0         2         0         0         7         1        55         0        20         0  
...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...   ...  
4114   12     0        1         1         0         2         2         0         3         2        50         0        20         0  
4115   21     0        1         1         3         0         2         0         1         3         0        216         0        20         0  
4116    9     8        2         2         3         0         0         0         0         6         1        61         1        20         1  
4117   40     0        1         1         3         0         0         0         0         1         0        510         0        20         0  
4118   16     4        2         2         3         0         2         0         0         7         4        172         0        20         0
```

4119 rows × 18 columns

```
In [21]: df_encoded['deposit'].value_counts()
```

```
Out[21]: deposit  
0    3668  
1     451  
Name: count, dtype: int64
```

```
In [22]: x = df_encoded.drop('deposit',axis=1) # independent variable  
y = df_encoded['deposit'] # dependent variable  
print(x.shape)  
print(y.shape)  
print(type(x))  
print(type(y))  
  
(4119, 17)  
(4119,)  
<class 'pandas.core.frame.DataFrame'>  
<class 'pandas.core.series.Series'>
```

```
In [23]: from sklearn.model_selection import train_test_split
```

```
In [24]: x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.25,random_state=1)  
print(x_train.shape)  
print(x_test.shape)  
print(y_train.shape)  
print(y_test.shape)  
  
(3089, 17)  
(1030, 17)  
(3089,)  
(1030,)
```

```
In [25]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
```

```
def eval_model(y_test,y_pred):  
    acc = accuracy_score(y_test,y_pred)
```

```

print('Accuracy_Score',acc)
cm = confusion_matrix(y_test,y_pred)
print('Confusion Matrix\n',cm)
print('Classification Report\n',classification_report(y_test,y_pred))

def mscore(model):
    train_score = model.score(x_train,y_train)
    test_score = model.score(x_test,y_test)
    print('Training Score',train_score)
    print('Testing Score',test_score)

```

In [26]: `from sklearn.tree import DecisionTreeClassifier`

```

dt = DecisionTreeClassifier(criterion='gini',max_depth=5,min_samples_split=10)
dt.fit(x_train,y_train)

```

Out[26]: `DecisionTreeClassifier` ⓘ ⓘ

► Parameters

In [27]: `mscore(dt)`

```

Training Score 0.923276141146002
Testing Score 0.9116504854368932

```

In [28]: `ypred_dt = dt.predict(x_test)`  
`print(ypred_dt)`

```
[0 0 1 ... 0 0 0]
```

In [29]: `eval_model(y_test,ypred_dt)`

```

Accuracy_Score 0.9116504854368932
Confusion Matrix
[[913 17]
 [ 74 26]]
Classification Report
      precision    recall   f1-score   support
          0       0.93     0.98     0.95      930
          1       0.60     0.26     0.36      100
      accuracy         0.91     0.91      1030
      macro avg       0.76     0.62     0.66      1030
  weighted avg       0.89     0.91     0.90      1030

```

In [30]: `from sklearn.tree import plot_tree`

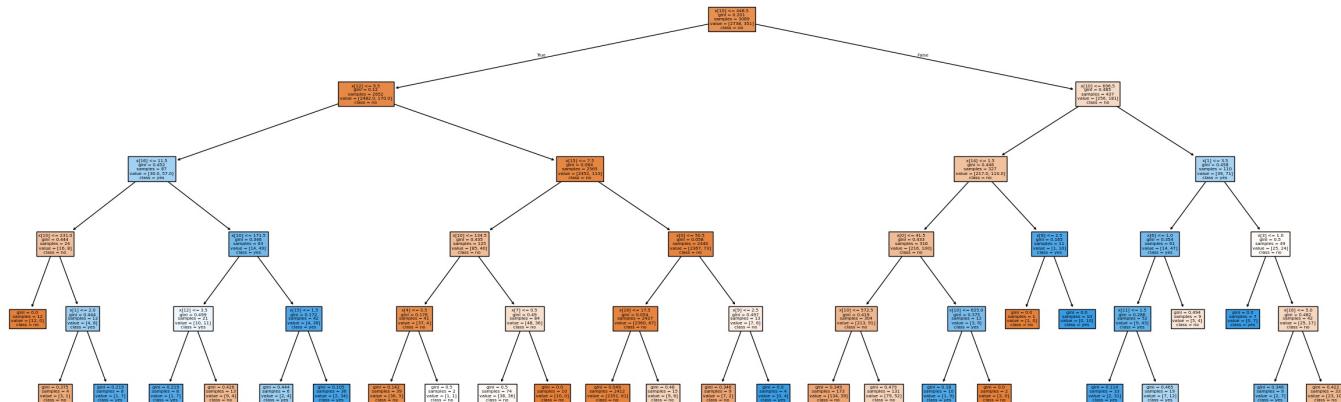
In [31]: `cn = ['no','yes']`  
`fn = x_train.columns`  
`print(fn)`  
`print(cn)`

```

Index(['age', 'job', 'marital', 'education', 'default', 'housing', 'loan',
       'contact', 'month', 'day_of_week', 'duration', 'campaign', 'pdays',
       'previous', 'poutcome', 'cons.price.idx', 'cons.conf.idx'],
      dtype='object')
['no', 'yes']

```

In [32]: `plt.figure(figsize=(30,10))`  
`plot_tree(dt,class_names=cn,filled=True)`  
`plt.show()`



In [33]: `dt1 = DecisionTreeClassifier(criterion='entropy',max_depth=4,min_samples_split=15)`

```
dt1.fit(x_train,y_train)
```

Out[33]: ▾ DecisionTreeClassifier ⓘ ?

► Parameters

In [34]: `mscore(dt1)`

Training Score 0.9145354483651668  
Testing Score 0.916504854368932

In [35]: `ypred_dt1 = dt1.predict(x_test)`

In [36]: `eval_model(y_test,ypred_dt1)`

Accuracy\_Score 0.916504854368932

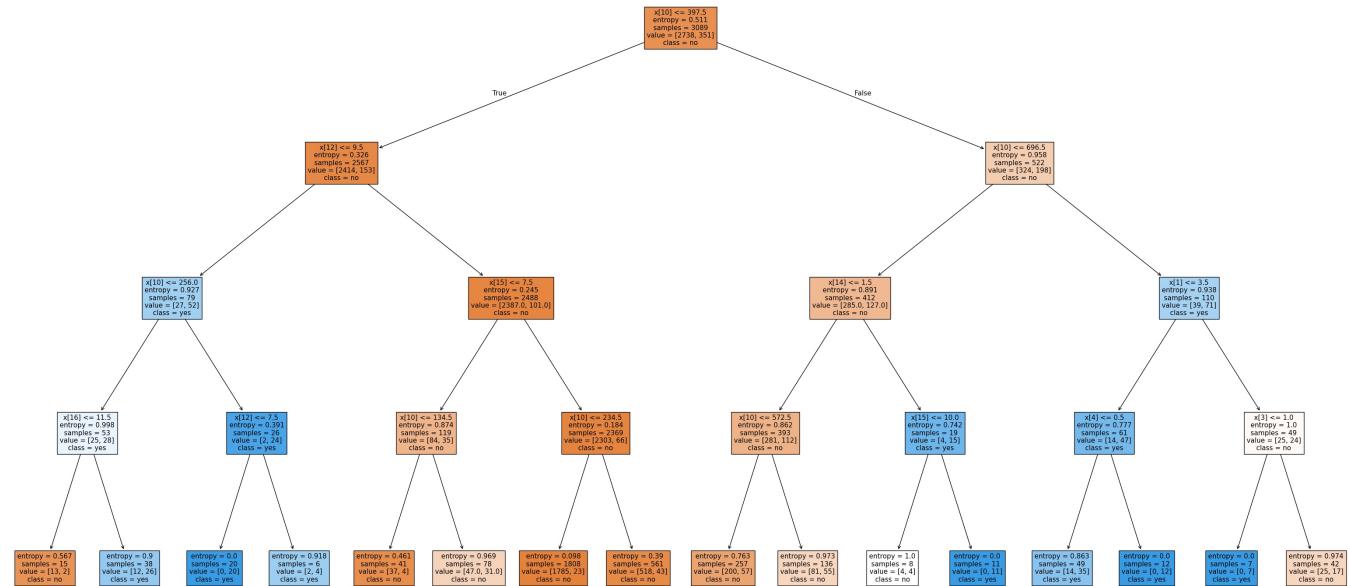
Confusion Matrix

```
[[912 18]
 [ 68 32]]
```

Classification Report

	precision	recall	f1-score	support
0	0.93	0.98	0.95	930
1	0.64	0.32	0.43	100
accuracy			0.92	1030
macro avg	0.79	0.65	0.69	1030
weighted avg	0.90	0.92	0.90	1030

In [37]: `plt.figure(figsize=(40,20))
plot_tree(dt1,class_names=cn,filled=True)
plt.show()`



In [ ]: